

20171056_Assignment3_up

March 17, 2020

```
In [1]: import cv2
import sklearn.mixture
import matplotlib.pyplot as plt
import numpy as np
import igraph
```

1 Grabcut assignment

1.1 Grabcut class defined here

1.1.1 Implementation details

- The code has classes
 - class 0 - background
 - class 1 - probably background
 - class 2 - probably foreground
 - class 3 - foreground
- First select the bounding box which becomes the probable foreground class.
- Learn two GMM models
 - Foreground (class 2 and class 3)
 - Background (class 0 and class 1)
- Depending on the class labels edges from source, sink and neighboring nodes are weighted.
- We first construct the graph with number of image pixels + 2 nodes (source and sink).
- We know the probable foreground nodes and background nodes in the first iteration.
- Edge weights assigned given class
 - class 0
 - * Sink - pixel - K
 - * Source - pixel - 0
 - * pixel - pixel - 0 (if class 0 or 1)
 - * pixel - pixel - neighboring penalty

- class 1
 - * Sink - pixel - negative log likelihood (foreground)
 - * Source - pixel - negative log likelihood (background)
 - * pixel - pixel - 0 (if class 0 or 1)
 - * pixel - pixel - neighboring penalty
- class 2
 - * Sink - pixel - negative log likelihood (foreground)
 - * Source - pixel - negative log likelihood (background)
 - * pixel - pixel - 0 (if class 3 or 2)
 - * pixel - pixel - neighboring penalty
- class 3
 - * Sink - pixel - 0
 - * Source - pixel - K
 - * pixel - pixel - 0 (if class 2 or 3)
 - * pixel - pixel - neighboring penalty

- Once the graph is constructed with edge weights we perform source terminal mincut. This partitions the graph into 2 parts.
- The above is repeated until convergence.

```
In [5]: class GrabCut:
        def __init__(self, image, gamma = 50, option = 0, num_components = 5, conv_threshold

            '''
            Class constructor
            Input:
                image - H x W x 3 - input image
                conv_threshold - scalar - convergence threshold for grabcut
                bbox - list - bounding box coordinates
                gamma - scalar - gamma value for grabcut
                option - scalar
                    0 - 4 neighborhood
                    1 - 8 neighborhood
                num_components - scalar - number of gaussian components
                disp_output - scalar - 1) shows intermediate outputs
                                   0) does not show intermediate results

            Returns:
                None
            '''
            self.threshold = conv_threshold
            self.disp = disp_output
            self.image = image
            if bbox is None:
                r = cv2.selectROI(self.image)
                bbox = [r[0], r[1], r[0] + r[2], r[1] + r[3]]
            self.bbox = bbox
```

```

else:
    self.bbox = bbox

self.option = option
self.x1 = int(bbox[0])
self.x2 = int(bbox[2])
self.y1 = int(bbox[1])
self.y2 = int(bbox[3])

self.gamma = gamma
self.mask = None
self.graph = None
self.num_components = num_components

self.bg_gmm = None
self.fg_gmm = None

self.create_mask_init()
self.create_graph_image()
cv2.destroyAllWindows()

def create_mask_init(self):
    '''
    Creates initial mask
    0 - bg
    1 - probably bg
    2 - probably fg
    3 - fg
    '''

    mask = np.zeros((self.image.shape[0], self.image.shape[1]))
    mask[ self.y1:self.y2, self.x1:self.x2] = 2
    self.mask = mask

def GMPrediction(self):
    '''
    Function to cluster images sections using gaussian mixture models
    '''

    image = self.image
    mask = self.mask

    bg_mask = np.logical_or(mask == 0 , mask == 1)
    fg_mask = np.logical_or(mask == 2 , mask == 3)

    image_mask_vector_bg = image[bg_mask, :]
    image_mask_vector_fg = image[fg_mask, :]

```

```

gmm_model_bg = sklearn.mixture.GaussianMixture(n_components = self.num_component
gmm_model_fg = sklearn.mixture.GaussianMixture(n_components = self.num_component

image_vector = image.reshape((-1, 3))

neg_log_likelihood_bg = -gmm_model_bg.score_samples(image_vector)
neg_log_likelihood_fg = -gmm_model_fg.score_samples(image_vector)

self.bg_gmm = neg_log_likelihood_bg.reshape((image.shape[0], image.shape[1]))
self.fg_gmm = neg_log_likelihood_fg.reshape((image.shape[0], image.shape[1]))

def create_graph_image(self):
    '''
    Function to create graph of image
    '''

    self.graph = igraph.Graph()
    self.graph.add_vertices(self.image.shape[0] * self.image.shape[1] + 2)

def calc_beta_value(self):
    '''
    Function to calculate beta value used in neighboring edge weight
    '''

    image = self.image
    height = image.shape[0]
    width = image.shape[1]

    diff_up = np.sum(np.sum(np.square(image[1:height - 1, :] - image[:height - 2, :])
    diff_side = np.sum(np.sum(np.square(image[:, 1:width - 1] - image[:, :width - 2])
    diff_diag = np.sum(np.sum(np.square(image[1:height - 1, 1:width - 1] - image[:height
    diff_lower_diag = np.sum(np.sum(np.square(image[:height - 2, 1:width - 1] - image

    #  $2 * (h - 1)(w - 1) + (h - 1)w + (w - 1)h = 4hw - 3h - 3w + 2$ 
    beta = (1 / 2) * (2 * (height - 1) * (width - 1) + height * (width - 1) + (height

    return beta

def add_graph_edges(self):
    '''
    Function to add edges and their respective weights to the graph
    '''

    beta = self.calc_beta_value()
    image = self.image

```

```

height = image.shape[0]
width = image.shape[1]

sink_node = int(image.shape[0]*image.shape[1] + 1)
source_node = int(0)

gamma = self.gamma
edges = []
weights = []
mask = self.mask

for row in range(height):
    for column in range(width):

        # bottom pixel weight

        if (row + 1) < height:
            edges.append((row*width + column + 1, (row + 1)*width + column + 1))

            if mask[row + 1, column] == mask[row, column] or (mask[row + 1, column] < mask[row, column]):
                weights.append(0)
            else:
                weights.append(gamma*np.exp(-beta * np.sum(np.square(image[row + 1, column] - image[row, column]))))

        #right pixel

        if (column + 1) < width:
            edges.append((row*width + column + 1, (row)*width + column + 2))

            if mask[row, column + 1] == mask[row, column] or (mask[row, column + 1] < mask[row, column]):
                weights.append(0)
            else:
                weights.append(gamma*np.exp(-beta * np.sum(np.square(image[row, column + 1] - image[row, column]))))

        # diagonal pixel

        if (column + 1) < width and (row + 1) < height and self.option == 1:
            edges.append(((row + 1)*width + column + 2, (row)*width + column + 1))

            if mask[row + 1, column + 1] == mask[row, column] or (mask[row + 1, column + 1] < mask[row, column]):
                weights.append(0)
            else:
                weights.append(gamma*np.exp(-beta * np.sum(np.square(image[row + 1, column + 1] - image[row, column]))))

K = 1 + max(weights)

for row in range(height):

```

```

    for column in range(width):

        if mask[row, column] == 1 or mask[row, column] == 2:
            edges.append((source_node, int(row*width + column + 1)))
            weights.append(self.bg_gmm[ row, column])

            edges.append((sink_node, int(row*width + column + 1)))
            weights.append(self.fg_gmm[ row, column])

        elif mask[row, column] == 0:
            edges.append((source_node, int(row*width + column + 1)))
            weights.append(0)

            edges.append((sink_node, int(row*width + column + 1)))
            weights.append(K)
        else:
            edges.append((source_node, int(row*width + column + 1)))
            weights.append(K)

            edges.append((sink_node, int(row*width + column + 1)))
            weights.append(0)

    self.graph.add_edges(edges)
    self.graph.es['weight'] = weights

def update_mask(self, partition_source):
    """
    Update mask after st mincut
    """
    init_mask = self.mask
    mask = np.zeros((self.image.shape[0], self.image.shape[1]))
    mask[ self.y1:self.y2, self.x1:self.x2] = 1
    mask2 = mask.copy()
    mask2[self.y1:self.y2, self.x1:self.x2] = 1
    bbox_width = image.shape[1]

    # removing source node from partition
    partition_source = np.array(partition_source, dtype = int)

    partition_source = partition_source - 1
    mask_partition = partition_source >= 0

    partition_source = partition_source[mask_partition]

```

```

fg_row = partition_source // bbox_width
fg_col = partition_source % bbox_width

mask[fg_row, fg_col] = 2
mask[np.where(init_mask) == 3] = 3
self.mask = mask * mask2

def generate_image_mask(self):
    '''
    Generate segmented image
    '''

    masked_image = np.zeros(self.image.shape)
    mask = np.logical_or(self.mask == 2, self.mask == 3) * 1
    for i in range(3):
        masked_image[:, :, i] = self.image[:, :, i] * mask
    masked_image = masked_image.astype('uint8')

    return masked_image

def grabCut(self):
    '''
    Function performs grabcut
    Input:
        Class object
    Returns:
        masked_image - H x W x 3 - image containing foreground only
    '''

    running = True
    i = 0
    while running:
        self.GMMPrediction()
        self.add_graph_edges()
        mc = self.graph.st_mincut( int(0), image.shape[0]*image.shape[1] + 1, self.g

        if i != 0:
            threshold = (mc.value - E_init) / mc.value
            print(threshold)
            if threshold < self.threshold:
                running = False

        E_init = mc.value
        partition_source, partition_terminal = mc.partition
        self.update_mask(partition_source)
        i = i + 1
        masked_image = self.generate_image_mask()

```

```

masked_image = cv2.cvtColor(masked_image, cv2.COLOR_BGR2RGB)
if self.disp == 1:
    plt.imshow(masked_image)
    plt.show()

return masked_image


```

```
In [6]: if __name__ == "__main__":
```


```

    name = 'flower'
    image_name = './data/images/%s.jpg' % (name)
    bbox_name = './data/bboxes/%s.txt' % (name)
    bbox = np.loadtxt(bbox_name, delimiter = ' ')
    image = cv2.imread(image_name)
    gamma = 50
    option = 0
    num_components = 5
    convergence_threshold = 0.1
    print_output = 1
    grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
    image_output = grabcut_object.grabCut()


```




0.4846690730081838




0.3314579455319971




0.25023953165851426




0.19948358511151848




0.16660283583520244




0.14276839702053631




0.12502269705973068




0.11104850966726496



0.10003330856192441



0.0898919100540024




1.1.2 Testing with 8 neighbors


Inference

- With 8 neighbors the foreground segmentation is more uniform diagonally


```
In [7]: name = 'flower'
image_name = './data/images/%s.jpg' % (name)
bbox_name = './data/bboxes/%s.txt' % (name)
bbox = np.loadtxt(bbox_name, delimiter = ' ')
image = cv2.imread(image_name)
gamma = 50
option = 1
num_components = 5
convergence_threshold = 0.1
print_output = 1
grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
image_output = grabcut_object.grabCut()
```




0.48784311639911937




0.3276402273521335




0.25019220244636836




0.19809750603803203




0.1682033655744191




0.14034618541618352




0.12733181461437545




0.10982823742887424



0.10051433427503613



0.09031618583009131




1.1.3 Testing with higher value of gamma

Inference

- Giving high gamma value means we are stressing over the neighbors to belong to same class this

```
In [8]: name = 'flower'
        image_name = './data/images/%s.jpg' % (name)
        bbox_name = './data/bboxes/%s.txt' % (name)
        bbox = np.loadtxt(bbox_name, delimiter = ' ')
        image = cv2.imread(image_name)
        gamma = 100
        option = 0
        num_components = 5
        convergence_threshold = 0.1
        print_output = 0
        grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
        image_output = grabcut_object.grabCut()
        plt.imshow(image_output)
        plt.show()
```

```
0.48722826367366456
0.3291582149590456
0.25005990914366905
0.19816484687373787
0.16708732594467054
0.14145399023205224
0.12576589261663812
0.11056142690712724
0.10080799818767916
0.09000614041396057
```




1.1.4 Varying GMM components

GMM component = 1

```
In [9]: name = 'banana1'
        image_name = './data/images/%s.jpg' % (name)
        bbox_name = './data/bboxes/%s.txt' % (name)
        bbox = np.loadtxt(bbox_name, delimiter = ' ')
        image = cv2.imread(image_name)
        gamma = 50
        option = 0
        num_components = 1
        convergence_threshold = 0.1
        print_output = 0
        grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
        image_output = grabcut_object.grabCut()
        plt.imshow(image_output)
        plt.show()
```

```
0.48004695255834995
0.2984890645589716
0.24272279553262927
0.20011467122571516
0.16666146315652355
0.1429198643140213
0.12496023499405283
```

0.11115309569354133
0.0999602300614365




GMM components = 5

```
In [10]: name = 'banana1'
         image_name = './data/images/%s.jpg' % (name)
         bbox_name = './data/bboxes/%s.txt' % (name)
         bbox = np.loadtxt(bbox_name, delimiter = ' ')
         image = cv2.imread(image_name)
         gamma = 50
         option = 0
         num_components = 5
         convergence_threshold = 0.1
         print_output = 0
         grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
         image_output = grabcut_object.grabCut()
         plt.imshow(image_output)
         plt.show()
```

0.4835244309625366
0.32738216641414775
0.24934476851369308
0.19871776637329905
0.1666412551857611

```
0.14235503031675814
0.1254341325165165
0.11065429747657508
0.10010883892281666
0.09066831940429741
```



1.1.5 Tight bounding box

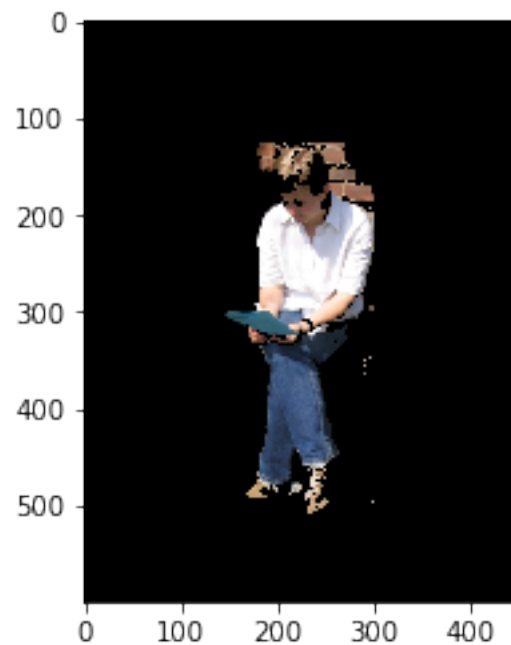
- Tighter bounding box gives higher accuracy
- With loose bounding box we have more pixels that are not in the foreground as probable foreground

```
In [11]: name = 'person6'
image_name = './data/images/%s.jpg' % (name)
bbox_name = './data/bboxes/%s.txt' % (name)
bbox = np.loadtxt(bbox_name, delimiter = ' ')
image = cv2.imread(image_name)
gamma = 50
option = 1
num_components = 1
convergence_threshold = 0.1
print_output = 0
grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
image_output = grabcut_object.grabCut()
plt.imshow(image_output)
plt.show()
```

```

0.49041299122697846
0.33180842954919404
0.2508846387920136
0.19893331409627926
0.16772006350435364
0.14175244450743651
0.12612353287380462
0.10996793516867336
0.10115597896841295
0.08973993369523252

```



1.1.6 Loose bounding box

```


In [12]: name = 'person6'
         image_name = './data/images/%s.jpg' % (name)
         bbox_name = './data/bboxes/%s.txt' % (name)
         bbox = np.loadtxt(bbox_name, delimiter = ' ')
         image = cv2.imread(image_name)
         gamma = 50
         option = 1
         num_components = 1
         convergence_threshold = 0.1
         print_output = 0
         grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
         image_output = grabcut_object.grabCut()

```



```
plt.imshow(image_output)
plt.show()
```

```
0.47755721295967557
0.33198249241625827
0.2503626144980008
0.1995347047184657
0.16710279546992302
0.14240116942188277
0.12546287591716515
0.11063996167209182
0.10047678609044333
0.09042723352747796
```



```
In [13]: name = 'cross'
image_name = './data/images/%s.jpg' % (name)
bbox_name = './data/bboxes/%s.txt' % (name)
bbox = np.loadtxt(bbox_name, delimiter = ' ')
image = cv2.imread(image_name)
gamma = 50
option = 0
num_components = 3
convergence_threshold = 0.1
print_output = 0
grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshol
```

```


image_output = grabcut_object.grabCut()
plt.imshow(image_output)
plt.show()

```

```

0.4848829345713043
0.33051042938323094
0.24910948967922233
0.19954943994714694
0.16666297824197926
0.14254639411659045
0.12506033325816562
0.11079659538476527
0.10007594768028226
0.09047683591795898

```



```

In [14]: name = 'cross'
          image_name = './data/images/%s.jpg' % (name)
          bbox_name = './data/bboxes/%s.txt' % (name)
          bbox = np.loadtxt(bbox_name, delimiter = ' ')
          image = cv2.imread(image_name)
          gamma = 50
          option = 0
          num_components = 5
          convergence_threshold = 0.1
          print_output = 0

```

```


grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshol
image_output = grabcut_object.grabCut()
plt.imshow(image_output)
plt.show()

```

```

0.4939997990455215
0.3283214037183475
0.2508581012445524
0.19707577983056299
0.16832506790367133
0.14089995889715876
0.12613626422670798
0.10967879541946883
0.10091810145696103
0.08960723419817657

```



1.1.7 testing with my own image

```

In [15]: name = 'chair_image'
image_name = './data/images/%s.jpg' % (name)
#bbox_name = './data/bboxes/%s.txt' % (name)
bbox = np.loadtxt(bbox_name, delimiter = ' ')
image = cv2.imread(image_name)
gamma = 50
option = 0

```

```


num_components = 3
convergence_threshold = 0.1
print_output = 0
grabcut_object = GrabCut(image, gamma, option, num_components, convergence_threshold)
image_output = grabcut_object.grabCut()
plt.imshow(image_output)
plt.show()

```

```

0.48405088351616776
0.3295249214759268
0.2503350275430667
0.19816026844742068
0.16771554486540816
0.14128109328902294
0.12618119616672308
0.10963334795923246
0.10123894504915058
0.08944276242329999

```



In []: